

1 Chapter 1

1.1 Exercises

1.1.1 Question 1.2

- (a) Absolute error ≈ 0.141592653 , the relative error ≈ 4.507 %.
- (b) Absolute error ≈ 0.001592653 , the relative error ≈ 0.0506957 %.
- (c) Absolute error ≈ 0.001264489 , relative error ≈ 0.0402499 %.

1.1.2 Question 1.5

The condition number is defined as follows:

$$Cond = \frac{|(f(\hat{x}) - f(x))/f(x)|}{|(\hat{x} - x)/x|}$$

but more carefully:

$$\begin{aligned} Cond &= \frac{|((x + \Delta x) - (y + \Delta y)) - (x - y)|/|x - y|}{|(|x + \Delta x| + |y + \Delta y|) - (|x| + |y|)|/(|x| + |y|)} \\ &\geq \frac{|((x - y) + \Delta x - \Delta y) - \epsilon|/|\epsilon|}{|(|x| + |\Delta x| + |y| + |\Delta y|) - (|x| + |y|)|/(|x| + |y|)} \\ &= \frac{|\Delta x - \Delta y|/\epsilon}{(|\Delta x| + |\Delta y|)/1} \\ &\geq \frac{1}{\epsilon} \end{aligned}$$

The condition number is therefore necessarily larger than $\frac{1}{\epsilon}$, which means as ϵ gets closer to 0, we have more sensitivity in f . Since $x - y \approx \epsilon$, this means that the function gets more sensitive the closer x and y are to each other in value.

1.1.3 Question 1.6

Forward error is $\hat{f}(x) - f(x)$, and the backward error is $\hat{x} - x$ where $f(\hat{x}) = \hat{f}(x)$.

(a)

$x=0.1$: The forward error is

$$0.1 - \sin(0.1) = 1.665 * 10^{-4},$$

the backward error is

$$\hat{x} - x = \sin^{-1}(0.1) - 0.1 = 0.100167 - 0.1 = 1.67 * 10^{-4}.$$

$x=0.5$: The forward error is

$$0.5 - \sin(0.5) = 0.0205,$$

the backward error is

$$\hat{x} - x = \sin^{-1}(0.5) - 0.5 = 2.36 * 10^{-2}.$$

$x=1.0$: The forward error is

$$1 - \sin(1) = 0.1585,$$

and the backward error is

$$\hat{x} - x = \sin^{-1}(1) - 1 = 0.5708.$$

(b) In this case, the forward error is

$$\hat{f}(x) - f(x) = x - \frac{x^3}{6} - \sin(x),$$

and the backward error can be found using the following:

$$backerr = \hat{x} - x,$$

where

$$\hat{x} = f^{-1}(\hat{f}(x)) = \arcsin(x - \frac{x^3}{6}),$$

$x=0.1$: The forward error is

$$0.1 - \frac{(0.1)^3}{6} - \sin(0.1) = -8.331 * 10^{-8},$$

and the backward error is

$$x - \hat{x} = 0.1 - \arcsin(0.1 - \frac{0.1^3}{6}) = 0.000000083731.$$

$x=0.5$: The forward error is

$$0.5 - \frac{(0.5)^3}{6} - \sin(0.5) = -2.588 * 10^{-4},$$

and the backward error is

$$x - \hat{x} = 0.5 - \arcsin(0.5 - \frac{0.5^3}{6}) = 0.00029495.$$

$x=1.0$: The forward error is

$$1.0 - \frac{(1.0)^3}{6} - \sin(1.0) = -8.137 * 10^{-3},$$

and the backward error is

$$x - \hat{x} = 1.0 - \arcsin(1.0 - \frac{1.0^3}{6}) = 0.0148892.$$

1.1.4 Question 1.10

(a) For x values near 0.

(b) The following is more stable around $x = 0$ because any rounding errors only occur once.

$$\frac{1}{1-x} - \frac{1}{1+x} = \frac{2x}{1-x^2}$$

1.1.5 Question 1.22

Assume a normalized floating-point system with $\beta = 10$, $t = 3$, and $L = -98$.

(a)

$$UFL = \beta^l = 10^{-98}$$

(b) If $x = 6.8710^{-97}$ and $y = 6.8110^{-97}$, what is the result of $x - y$?

Ignoring the system, $x - y = 0.06 * 10^{-97}$ which if we normalize (which in the initial case we must), we have the result $= 6 * 10^{-99}$ which is smaller than the underflow value. If we're assuming truncation, then we end up with a value of 0 for this result.

(c) What would be the result of $x - y$ if the system permitted gradual underflow?

If the system permitted gradual underflow, then the result would be

$$result = 0.60 * 10^{-98}.$$

1.2 Computer Exercises

1.2.1 Question 1.6

(a) The method which uses multiplication, ie

$$x_k = a + kh, \quad k = 0, \dots, n$$

is better than the recursive summation method. Since we are using floating point arithmetic, every step of the recursive method will involve some rounding. This loss of precision (error) will accumulate in the sum. While the multiplicative method also involves rounding, there are only two operations, so the loss of precision is restricted to whatever occurs because of those two operations. Basically there are two rounding errors for the multiplicative method while the recursive method adds n rounding errors.

(b) In order to do this we set $a = 0$, $b = 1$ and use $n = 10000$. We computed the x_k for both methods and compared the results by taking the squared error between each term of each method and a true set of values achieved using `np.linspace`. The code can be found in *assignment1.ipynb*.




Figure 1: Comparison of the error of the two proposed linspace generation methods.

1.2.2 Question 1.7




Figure 2: Plot of the errors for the finite difference and centered difference approximations of the derivative of $\sin(x)$ at $x = 1.0$ for different values of h .

There isn't really much to comment about for this question. The plot requested can be found in figure 1.2.2. Apart from that you can find the code in *assignment1.ipynb*.

1.2.3 Question 1.9

- (a) The code for this question can be found in *assignment1.ipynb*.
- (b) The natural way to stop is when the term we are adding hits the value of 0. This is because factorials grow faster than their equivalent polynomials, so at one point the denominator will be so much bigger than the numerator that it will round to zero in the computer.
- (c) The results are shown in figure 1.2.3

Relative error of infinite series approximation of $\exp(x)$ at various x values.




Figure 3: Plot of the relative error of the infinite series approximation of the exponential function.

1.2.4 Question 1.10

The code for this question can be found in assignment1.ipynb.

The explanations for implementation are commented into the code. Effectively, the first step was to check for a linear equation, that is, if the a coefficient is effectively zero. If this is the case, then we have a linear equation. If the b coefficient is non-zero then we have a single root. If the b coefficient is also zero, then our function is a constant and it has no roots.

The next thing to check is to see if there are any imaginary roots. If there are we don't handle them, because we are told not to in the question spec.

Next we check if the square root of the discriminant is negligible (i.e. smaller than the machine epsilon). If it is then we have a repeating root case.

Finally, if all of those pass, then we need to select the standard and alternate methods for each root in order to avoid cancellation in the b coefficient and the determinant. This involves picking the formula which matches the sign for the two for each root.

Any further details can be found in the code. The results for the test cases are given below:

a	b	c	1 st root	2 nd root
6	5	-4	0.5	-1.3333333333333333
6e+30	5e+30	-4e+30	0.5	-1.3333333333333335
0	1	1	-1.0	nan
1	-100000.0	1	99999.99999	1.0000000001000001e-05
1	-4	3.99999999	2.0000999999999696	1.9999000000000304
1e-30	-1e+30	1e+30	1.0	nan

1.2.5 Question 1.14

The code for this question can be found in assignment1.ipynb.




Figure 4: The value of the different representations of $(x - 1)^6$ plotted over the range $[0.995, 1.005]$.

There are multiple reasons why this issue is occurring in this circumstance. The first is that there are more opportunities for cancelation, because there are more subtractions in the function. The second, and probably the more important one, is that for values so close to one, the higher powers in the expanded terms are driven very close to 1, so the cancelation that happens in these terms is much more severe. In the compact term, the cancelation is minimal, because the subtraction occurs when the values have not been pulled close together, so the cancelation isn't bad. Then with this relatively unperturbed value, we take the sixth power, which doesn't cause much error.

2 Chapter 2

2.1 Exercises

2.1.1 Question 2.7

We have the matrix $A = \begin{pmatrix} 1 & 1+\epsilon \\ 1-\epsilon & 1 \end{pmatrix}$

(a) The determinant of this matrix is

$$\begin{aligned} & 1 - (1 + \epsilon) * (1 - \epsilon) \\ &= 1 - 1 + \epsilon - \epsilon + \epsilon^2 \\ &= \epsilon^2 \end{aligned}$$

(b) The value of epsilon is the value of the machine's epsilon. In floating point, the value of ϵ must be less than the square root of the machine epsilon value. More formally the required condition is

$$\epsilon < \sqrt{\epsilon_{machine}}$$

(c) In order to get A into upper triangular, we need to annihilate the first entry of the second row, which we do by subtracting $(1 - \epsilon) * \text{the first row from the second row.}$

$$A = \begin{pmatrix} 1 & 0 \\ -(1-\epsilon) & 1 \end{pmatrix} * \begin{pmatrix} 1 & 1+\epsilon \\ 1-\epsilon & 1 \end{pmatrix} = \begin{pmatrix} 1 & 1+\epsilon \\ 0 & \epsilon^2 \end{pmatrix} = LU$$

so $L = \begin{pmatrix} 1 & 1+\epsilon \\ 1-\epsilon & 1 \end{pmatrix}$ and $U = \begin{pmatrix} 1 & 1+\epsilon \\ 0 & \epsilon^2 \end{pmatrix}$.

(d) The determinant of A and U are the same, so the same thing must hold in order for U to be singular (i.e. the determinant be 0).

$$\epsilon < \sqrt{\epsilon_{machine}}$$

2.1.2 Question 2.21

$$x = B^{-1}(2A + I)(C^{-1} + A)b$$

Let's go right to left. First thing is first, we can compute $x = Ab$, as well as solve out the inverted C matrix which is giving us trouble by setting $Cy = b$ and solving for y . We then have

$$x = B^{-1}(2A + I)(y + x)$$

let's combine x and y to give us z , so $z = x + y$.

$$x = B^{-1}(2A + I)z$$

The rightmost term can be called $w = (2A + I)z$ we now have

$$x = B^{-1}w.$$

Finally we solve the equation $Bx = w$ for x and we have our result for x without inverting any matrices.

2.1.3 Question 2.22

We are factorizing an \mathbf{n} by \mathbf{n} matrix. The reduction of the matrix into L and U by Gaussian Elimination requires, at each row i

1. The selection of the annihilation value for each of the $n - i$ rows below.
2. The multiplication of the current row by each of these factors and the subsequent subtraction of these scaled values from the below rows in the U matrix. This is the annihilation step.
3. The matrix multiplication of the current cumulative L by the annihilating operation matrix. This I'll call the accumulation step.

For the following complexity assume the matrices are indexed starting at 1.

[H] Step	Approximate Complexity (size n and current row i)
Selection	$n - i$ (only multiplications)
Annihilation	$(n - i + 1)(n - i)$ (both multiplications and subtractions)
Accumulation	I can't figure out the complexity.

The first step is negligible in terms of complexity. The second step is by far the dominant step in terms of complexity. We have effectively $(n - i + 1)(n - i)$ multiplications and subtractions at step i with a matrix of size n .

Over the whole process, we then have

$$\sum_{i=1}^{i=n} (n - i + 1)(n - i) \approx \sum_{i=1}^{i=n} (n - i + 1)^2 = \sum_{i=1}^{i=n} i^2 = \frac{n(n + 1)(2n + 1)}{6}$$

therefore the dominant term is the $\frac{2n^3}{6}$ term which reduces to

$$\frac{n^3}{3}.$$

2.1.4 Question 2.32

We need to show that the following functions satisfy the first three properties of a matrix norm.

The matrix norms we have defined satisfy the following important properties, where A and B are any matrices:

1. $\|A\| > 0$ if $A \neq O$.
2. $\|\gamma A\| = |\gamma| * \|A\|$ for any scalar γ .
3. $\|A + B\| \leq \|A\| + \|B\|$.

(a) The norm is given:

$$||A||_{max} = \max_{i,j} |a_{i,j}|$$

(i) If $A \neq 0$ then some $a_{ij} \neq 0$ so $|a_{ij}| > 0$ therefore

$$\max_{i,j} |a_{i,j}| > 0 \quad \text{if} \quad A \neq 0.$$

(ii) For the second condition:

$$||\gamma A||_{max} = \max_{i,j} |\gamma a_{ij}| = |\gamma| \max_{i,j} |a_{ij}| = |\gamma| * ||A||_{max}$$

(iii) Finally we have for the third condition, with matrices A and B let's assume we have the following:

$$||A||_{max} = |a_{i,j}|$$

$$||B||_{max} = |b_{k,l}|$$

The most aggressive case would be if $a_{i,j}$ and $b_{k,l}$ are the same sign. In this case we have

$$||A + B||_{max} = |a_{i,j} + b_{k,l}| = |a_{i,j}| + |b_{k,l}| = ||A||_{max} + ||B||_{max}$$

This is the equality case, and it is a true equality, that is scaling either of the max values leads to the same increase in the norm of the combination.

If the signs are flipped (i.e the elements which have the largest absolute value in both A and B are opposite sign) then we cannot possibly have equality, and we've already shown that we can't go above equality, so we must have a smaller value in the combined norm than in the sum of the two separate norms.

(b)

$$||A||_F = \left(\sum_{i,j} |a_{i,j}|^2 \right)^{\frac{1}{2}}$$

2.2 Computer Exercises

2.2.1 Question 2.5

The code can be found in assignment1.ipynb.

Part A

The value of x after single precision gaussian elimination is

$$\begin{pmatrix} -0.99992704 \\ 2.000323 \\ -2.9988637 \\ 3.9983127 \end{pmatrix}$$

Part B

The value of the residual r is

$$\begin{pmatrix} 4.7683716^{-7} \\ -1.0561943^{-4} \\ 2.0742416^{-5} \\ -9.5858813^{-6} \end{pmatrix}$$

Part C

The value of the correction z is

$$\begin{pmatrix} -7.2929193^{-5} \\ -3.2293634^{-04} \\ -1.1358784^{-03} \\ 1.6866578^{-03} \end{pmatrix}$$

The corrected value of x is

$$\begin{pmatrix} -1. \\ 2.0000002 \\ -2.9999995 \\ 3.9999993 \end{pmatrix}$$

Part D

The corrected value of x is (after 2 iterations)

$$\begin{pmatrix} -1. \\ 2. \\ -3. \\ 4. \end{pmatrix}$$

2.2.2 Question 2.6

Again, the code can be found in *assignment1.ipynb*.

The error hits 100% at $n = 13$. The number of significant digits has an inverse relationship with the log of the condition number. This makes sense since the significant digits is acquired by taking the log of the error and rounding to the nearest digit.

Significant digits and condition number for Hilbert matrix of size $n \times n$.




Figure 5: The number of significant digits plotted against the log of the condition number as we increase the size n of the Hilbert matrix.

2.2.3 Question 2.9

In both cases we do pretty poorly. The code can be found in *assignment1.ipynb*. The results are as follows:

ϵ	Sig Figs G.E.	Sig Figs r correction
0.01	inf	inf
0.0001	inf	inf
1e-06	7	7
1e-08	0	0
1e-10	0	0
1e-12	0	0
1e-14	0	0
1e-16	0	0
1e-18	0	0
1e-20	0	0

As we can see, the correction has no impact on the result, this is because the computation of the residual is subject to the same troubling truncation issues as the Gaussian Elimination solver is.